

Lab 3 Part 4 Report

Explain

The output of Explain has 12 columns: id, select_type, table, partitions, type, possible_keys, key, key_len, ref, rows, filtered, and Extra. These columns describe how the output of the query is produced.

The **id** describes the SELECT identifier.

The **select_type** describes the SELECT type.

The **table** describes the table for the output rows.

The **partitions** describes the matching partitions.

The **type** describes how it's joined.

The **possible_keys** describes what indexes can be chosen.

The **key** describes which index is actually chosen.

The **key_len** describes the length of this chosen key.

The **ref** describes the columns compared to the index.

The **rows** describes an estimate of the number of rows that will be examined.

The **filtered** describes the percentage of rows that are filtered.

The **Extra** describes any other additional information, such as describing how it's joined.

Part 4-1

Before

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	Review	NULL	ALL	NULL	NULL	NULL	NULL	1120345	100.00	Using where

Run	Times (s)	Average	Standard Deviation
1	0.76	0.77	0.006519202
2	0.773		
3	0.776		
4	0.774		
5	0.767		

After

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	Review	NULL	index	NULL	review_date_idx	6	NULL	1120345	100.00	Using where; Using index

Run	Times (s)	Average	Standard Deviation
1	0.242	0.2466	0.005683309
2	0.251		
3	0.25		
4	0.251		
5	0.239		

Improvement

After creating an index in the Review table for the date column, we see the average execution time drop from 0.77 seconds to 0.2466 seconds – an execution time reduction of 68%.

Alternatives

In this case, the only logical choice of column to index on is the date column, since we are simply returning the number of reviews written in a specific month. For this specific query, no other choice of indexing makes sense.

Part 4-2

Before

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	ud	NULL	ALL	NULL	NULL	NULL	NULL	765240	10.00	Using where
1	SIMPLE	r	NULL	ALL	NULL	NULL	NULL	NULL	1120345	10.00	Using where; Using join buffer (hash join)
1	SIMPLE	b	NULL	ALL	NULL	NULL	NULL	NULL	183269	100.00	Using where; Using join buffer (hash join)

Run	Times (s)	Average	Standard Deviation
1	1.416	1.4274	0.006730527
2	1.43		
3	1.427		
4	1.431		
5	1.433		

After

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	r	NULL	ref	review_user_id	review_user_id	90	const	13	100.00	NULL
1	SIMPLE	ud	NULL	ALL	NULL	NULL	NULL	NULL	765240	10.00	Using where; Using join buffer (hash join)
1	SIMPLE	b	NULL	ALL	NULL	NULL	NULL	NULL	183269	100.00	Using where; Using join buffer (hash join)

Run	Times (s)	Average	Standard Deviation
1	0.59	0.5982	0.00626099
2	0.6		
3	0.606		
4	0.601		
5	0.594		

Improvement

After creating an index in the Review table for the user_id columns, we see the average execution time drop from 1.4274 seconds to 0.5982 seconds – an execution time reduction of 58%.

Alternatives

Another option would be to have a single index on business_id in the Review table to be used join the Business table and the Review table on business_id more efficiently.